

Benchmarking Rule-Based, Learned, and Zero-Shot Vision-Language Grasp Prediction Under a Unified Evaluation Protocol

Anish Talla
High School Student
United States

Abstract—I compared three ways to predict a robotic grasp from an RGB image: a detector-assisted heuristic, a neural network trained on grasp labels, and a general-purpose vision-language model (VLM) used zero-shot. All three systems were tested on the same object-wise split of the Cornell Grasping Dataset and scored by the same rectangle metric. A fine-tuned ResNet18 passed on 79.7% of test images. The heuristic passed on a post-hoc-corrected 57.7%, up from 40.7% before the correction. GPT-4o passed on 12.4% of calls across five independent runs per image. Object-clustered resampling left the order unchanged. Their errors were different. The heuristic struggled with orientation, the ResNet18 usually missed on placement, and GPT-4o failed both criteria on 59.3% of calls. Its predicted center lay outside the convex hull containing all labeled grasp corners on 53.4% of parsed calls, compared with 7.0% and 6.5% for the other systems. The experiment covers one model and one prompt, so it cannot support a claim about VLMs in general. A follow-up should test whether a model that can name a reasonable grasp region loses accuracy when it must translate that choice into pixel coordinates.

Index Terms—robotic grasping, computer vision, deep learning, vision-language models, multimedia data mining, spatial grounding

1. Introduction

A robot has to make a geometric decision before it can close its gripper: where the fingers should land, how the gripper should rotate, and how far the jaws should open. A programmer can write that decision into a set of rules. A neural network can learn it from labeled examples. A VLM can also be asked for the answer directly, even if it has never been trained for grasp prediction. The cost and setup differ sharply across these approaches, yet their reported accuracies are hard to compare when the input, output, or scoring rule changes with the method.

I built one system of each kind and kept those parts fixed. Each system saw the same 123 held-out RGB images, returned the same geometric representation, and was scored against the same Cornell labels by one implementation of the standard rectangle test [1], [2]. I also reconstructed object identities so that no physical object crossed between training and test. The statistical analysis resamples whole objects

because several test images are rotated views of the same item.

I first ask which system predicts the most valid rectangles, then look at where each system goes wrong. The three systems missed different parts of the metric. I report those failures with a common taxonomy and a separate check of whether each predicted center lands near any labeled grasp region.

Current dense grasp detectors already report higher Cornell scores [4]. My aim here is different: compare three practical starting points without changing the test when the method changes.

2. Related Work

Jiang *et al.* introduced the Cornell dataset and its rectangle representation of a grasp [1]. The rectangle stores the center, orientation, opening, and jaw depth. Under the usual Cornell test, a prediction passes when it is within 30° of a labeled grasp and its intersection over union (IoU) is greater than 25% [2]. Redmon and Angelova used a neural network to regress one rectangle and reported 84.9% object-wise accuracy with RGB-D input [3]. GR-ConvNet and other pixel-wise models report higher scores, although they make a dense prediction instead of regressing one rectangle for the whole image [4].

VLM-based grasping systems usually do not ask the language model for the final pixel coordinates. Set-of-Mark prompting has the model choose among regions already marked on an image [5]. VLAD-Grasp has a VLM draw the grasp, then recovers a pose with depth prediction, segmentation, and point-cloud alignment [6]. System C tests a smaller pipeline: show GPT-4o the image and ask it to return the coordinates itself. Its score should therefore be read as the result of that output method, not as a limit on VLM-based grasping.

3. Experimental Design

3.1. Dataset, Split, and Metric

Cornell contains 885 RGB-D images of household objects, with several positive grasp rectangles on most images. I used only RGB. Two images were dropped because their object boundaries could not be resolved, leaving 883.

TABLE 1. OBJECT-WISE CORNELL SPLIT.

Partition	Objects	Images
Train	164	620
Validation	35	140
Test	35	123
Total	234	883

Cornell does not include object IDs, so I reconstructed them from the frame sequence. The procedure segmented the object from the photography platform and compared consecutive masks with color and area descriptors that tolerate rotation. Ambiguous boundaries were reviewed manually. I set the thresholds asymmetrically because the two possible mistakes have different costs. If two objects are merged, the combined group still stays in one split. If one object is split in two, its images can leak into both training and test.

An automated first pass assigned a decision and confidence to each ambiguous boundary. In a blind 30-boundary audit, its calls agreed with my labels on 86.7% of confident cases and 40.0% of uncertain ones. I then decided every uncertain call and every high-risk “different object” call. Table 1 shows the final seed-42 split. A second check confirmed that no object appears in more than one partition.

I scored each prediction against every positive rectangle on its image. A match to any one of them counts as correct. Before using the evaluator at scale, I checked hand-computed cases, rectangle round trips, and an independent rasterized IoU calculation. All predictions were converted back to the native 640×480 coordinate frame for the final comparison.

3.2. System A: Detector-Assisted Heuristic

System A starts with a COCO-pretrained Faster R-CNN using a ResNet50-FPN-v2 backbone [7]. A fixed lookup table maps the detected category to a center, upper, or end grasp. I calibrated the global opening and jaw-size constants on training labels only. The COCO vocabulary misses most Cornell objects, so a platform-segmentation fallback supplies a box when the detector cannot. The rule closes the gripper across the box’s narrow dimension. It can output only 0° or 90° .

The first test run scored 40.7%. Its rendered predictions exposed a defect: some accepted detections covered background clutter instead of the object on the platform. I added a guard requiring the detection to contain the segmented object’s centroid and recalibrated the global constants on training data. The score rose to 57.7%. Test output prompted that change, so 57.7% is post-hoc and 40.7% remains the clean pre-correction result.

3.3. System B: Learned Predictor

System B replaces ResNet18’s ImageNet classifier [8] with separate heads for center, orientation, and size. The orientation head predicts $(\cos 2\theta, \sin 2\theta)$. Doubling the angle

handles the fact that a parallel-jaw grasp is unchanged by a 180° rotation. For images with several labels, I compute the loss against every rectangle and backpropagate only the best match. Averaging all valid labels could put the target between two good grasps, where neither one actually exists.

I trained a custom CNN, ResNet18, and ResNet34 on the 620 training images. Affine augmentation moved the rectangle corners with the pixels. Architecture, checkpoint, size-loss weight, and the custom CNN’s learning rate were chosen on validation data. ResNet18 had the best validation accuracy at 83.6%, so I froze that model before opening the test results.

3.4. System C: Zero-Shot VLM

System C sends the native RGB image to GPT-4o through OpenRouter [9]. It uses no grasp-specific training. The prompt asks for two fingertip coordinates and a jaw width in JSON, which together define the rectangle. I developed the prompt on 30 training images and froze it before test.

Each test image was sent in five independent, single-turn calls at the provider’s default temperature. I retried transport failures, but not replies that arrived and failed to parse. Retrying those would quietly turn a single-call baseline into a best-of-several system. Of 615 calls, 614 parsed. The primary score pools the five exchangeable runs to estimate one-call accuracy. Best-of-five is only a ceiling: it requires five calls and an oracle that knows when one of them is right.

3.5. Statistical and Failure Analysis

The 123 test images are repeated views of 35 unseen objects, not 123 independent samples. I formed the 95% confidence intervals from 20,000 object-clustered bootstrap resamples. Every view of an object, along with all five VLM calls, moves together. Paired differences use the same resamples and 20,000 object-level label-swap permutations.

The failure analysis gives each attempt one label: correct, angle only, overlap only, both criteria, or no prediction. I also checked whether the predicted center falls inside the convex hull of all labeled grasp corners on the image. This is a loose spatial check, not a claim that the grasp would work on hardware. A center outside that hull falls beyond the area spanned by the annotated grasp corners.

4. Results

4.1. Overall Accuracy

ResNet18 passed 98 of 123 images (79.7%). The corrected heuristic passed 71 (57.7%), and GPT-4o passed 76 of 615 calls (12.4%). Its five runs ranged from 9.8% to 14.6%. At least one GPT-4o call passed on 43 of the 123 images, which gives a best-of-five ceiling of 35.0%. Table 2 includes the object-clustered intervals.

TABLE 2. ACCURACY WITH OBJECT-CLUSTERED 95% CONFIDENCE INTERVALS.

System	Accuracy	95% CI
A: heuristic, corrected	57.7%	[46.2, 68.2]
B: ResNet18	79.7%	[70.8, 87.1]
C: GPT-4o, one call	12.4%	[8.0, 17.4]
C: best of five	35.0%	[24.6, 45.5]

TABLE 3. SHARED FAILURE TAXONOMY. PERCENTAGES USE ALL ATTEMPTS.

Outcome	A ($n = 123$)	B ($n = 123$)	C ($n = 615$)
Correct	71 (57.7)	98 (79.7)	76 (12.4)
Angle only	13 (10.6)	4 (3.3)	85 (13.8)
Overlap only	18 (14.6)	15 (12.2)	88 (14.3)
Both	13 (10.6)	6 (4.9)	365 (59.3)
No prediction	8 (6.5)	0 (0.0)	1 (0.2)

ResNet18 led the corrected heuristic by 22.0 points. The bootstrap interval for that difference was [9.7, 35.1], with permutation $p = 0.0031$. The corrected heuristic led GPT-4o’s best-of-five ceiling by 22.8 points [9.2, 35.6], $p = 0.0060$, and its expected one-call score by 45.4 points [33.8, 55.6], $p < 0.0001$. Giving each object equal weight produced 51.3%, 75.4%, and 10.5% for A, B, and one-call C. That calculation changed the size of the gaps but not their order.

4.2. Failure Patterns

The error counts in Table 3 separate the systems more clearly than accuracy alone. Thirteen of A’s 52 failures had enough overlap and missed only the angle test, which fits a system limited to two orientations. B had four angle-only failures. Fifteen of its 25 misses failed only on overlap, so placement was the larger remaining problem. C failed both tests on 365 of 615 calls (59.3%).

The center-hull check shows where C’s failures begin. A landed outside the labeled grasp hull on 8 of 115 predictions (7.0%), and B did so on 8 of 123 (6.5%). C did so on 328 of 614 parsed calls (53.4%). Figure 1 shows the shared coordinate frame on one image. This is not an even contest by itself: A starts from a detected or segmented object, and B was trained on Cornell. C can place its coordinates anywhere. I use this check to locate a possible stage of failure. The three rates do not estimate the same underlying capability.

The five GPT-4o calls also disagreed with one another. Mean pairwise rectangle agreement was 22.0%, and mean pairwise IoU was 0.13. All five calls failed on 80 of 123 images. All five passed on one image. The remaining 42 images were unstable: identical pixels sometimes produced a passing grasp and sometimes did not.

5. Discussion and Limitations

On this setup, asking GPT-4o for raw coordinates worked poorly. This result does not extend to VLM grasping methods in general. Set-of-Mark and VLAD-Grasp use a VLM




Figure 1. One test image under the shared coordinate system. Labeled grasps are green, System A red, System B blue, and System C’s five repeats orange.

without making it emit the final coordinates, and both add machinery that System C does not have [5], [6].

On a sample of C’s failures, I read rationales that named a plausible part of the object while the coordinates landed elsewhere. That reading was informal, not a blinded coding study. The numerical result is the 53.4% center-hull rate. Together they suggest a *coordinate-binding hypothesis*: GPT-4o may recognize a useful region in words and then lose it while mapping the answer into pixels. A fair test would give the same model the same tuning budget, then compare coordinate output with marks generated without using the ground-truth grasps.

The experiment has several limits. It uses one object-wise split of a small benchmark, RGB without depth, and no robot trials. Cornell’s annotations are not exhaustive, so a rectangle that fails the metric might still work on a real object. System B also received more development effort: three architectures and validation-based tuning, compared with one frozen prompt for C. An early formatting check on about 550 grasp rectangles probably included a few images that later entered test, although no parameter or rule came from that check. Test inspection also prompted System A’s correction. Its 40.7% score is the clean pre-correction result; comparisons using 57.7% are post-hoc. Confirmation requires a fresh split, object-wise cross-validation, or a second dataset.

6. Conclusion

ResNet18 scored highest at 79.7%. The corrected heuristic reached 57.7%, and GPT-4o reached 12.4% per call. Their error locations add context to that ranking. The heuristic often found the object but could not represent its angle. ResNet18 usually got the angle and missed the placement. GPT-4o often failed before either distinction mattered because its rectangle was not near a labeled grasp region.

A direct next test would give the same VLM marked candidate regions instead of asking for coordinates. The candidates would have to be generated without the grasp labels, and the result would be scored on unseen objects. I have not run that experiment. It would help distinguish a visual-perception failure from a failure in turning a visual choice into numbers.

Acknowledgment

Anthropic Claude Code helped with code development, including implementation of the grasp detection model, and helped organize ideas and answer questions during the writing process. I designed the study, made the research decisions, ran and checked the experiments, wrote the manuscript, and verified the citations and numerical claims. GPT-4o was evaluated separately as System C in Section III-C.

References

- [1] Y. Jiang, S. Moseson, and A. Saxena, “Efficient grasping from RGBD images: Learning using a new rectangle representation,” in *Proc. IEEE Int. Conf. Robot. Autom. (ICRA)*, 2011, pp. 3304–3311.
- [2] I. Lenz, H. Lee, and A. Saxena, “Deep learning for detecting robotic grasps,” *Int. J. Robot. Res.*, vol. 34, nos. 4–5, pp. 705–724, 2015.
- [3] J. Redmon and A. Angelova, “Real-time grasp detection using convolutional neural networks,” in *Proc. IEEE Int. Conf. Robot. Autom. (ICRA)*, 2015, pp. 1316–1322.
- [4] S. Kumra, S. Joshi, and F. Sahin, “Antipodal robotic grasping using generative residual convolutional neural network,” in *Proc. IEEE/RSJ Int. Conf. Intell. Robots Syst. (IROS)*, 2020, pp. 9626–9633.
- [5] J. Yang *et al.*, “Set-of-Mark prompting unleashes extraordinary visual grounding in GPT-4V,” *arXiv:2310.11441*, 2023.
- [6] M. Kulshrestha *et al.*, “VLAD-Grasp: Zero-shot grasp detection via vision-language models,” *arXiv:2511.05791*, 2025.
- [7] S. Ren, K. He, R. Girshick, and J. Sun, “Faster R-CNN: Towards real-time object detection with region proposal networks,” in *Adv. Neural Inf. Process. Syst.*, vol. 28, 2015.
- [8] K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2016, pp. 770–778.
- [9] OpenAI, “GPT-4o system card,” *arXiv:2410.21276*, 2024.