

Three Ways to Miss a Grasp: A Same-Metric Comparison of Rule-Based, Learned, and Zero-Shot Vision-Language Grasp Prediction

Anish Talla

High School Student

Lightridge High School and Academies of Loudoun

Virginia, United States

anish.talla99@gmail.com

Preprint, September 2026. Code and data: <https://github.com/anishtalla27/Vision-Based-Grasping>

Abstract—Accuracy numbers for multimodal models are hard to compare when every method is tested in its own way. I compared three ways to predict a robotic grasp from an RGB image under one fixed protocol: a detector-assisted heuristic, a neural network trained on grasp labels, and a general-purpose vision-language model (VLM) used zero-shot. All three systems were tested on the same object-wise split of the Cornell Grasping Dataset and scored by the same rectangle metric. A fine-tuned ResNet18 passed on 79.7% of test images, and a second training round with a stable recipe and three seeds per architecture put the learned system at 84.0%. The heuristic passed on a post-hoc-corrected 57.7%, up from 40.7% before the correction. GPT-4o passed on 12.4% of 615 calls, five independent runs per image, with low agreement between runs on identical pixels. Because test images are repeated views of 35 objects, all intervals come from object-clustered resampling, which widened every interval and left the order unchanged. The three systems also failed in different ways. The heuristic struggled with orientation, the ResNet18 usually missed on placement, and GPT-4o failed both criteria on 59.3% of calls. Its predicted center lay outside the convex hull containing all labeled grasp corners on 53.4% of parsed calls, compared with 7.0% and 6.5% for the other systems. That pattern suggested the model could see a grasp but not express it as coordinates. A fourth run tested this directly by showing the same model label-free candidate grasps drawn on the image and asking it to choose one by number. On three candidate menus, including one with every mark drawn the same size, its accuracy stayed at the random floor of the menu, and it preferred a grasp closing along the object’s long axis regardless of mark size. The model’s weak point is choosing the grasp. Writing the coordinates was never the main problem.

Index Terms—robotic grasping, vision-language models, GPT-4o, Set-of-Mark prompting, Cornell Grasping Dataset, clustered bootstrap

1. Introduction

A robot has to make a geometric decision before it can close its gripper: where the fingers should land, how the

gripper should rotate, and how far the jaws should open. A programmer can write that decision into a set of rules. A neural network can learn it from labeled examples. A VLM can also be asked for the answer directly, even if it has never been trained for grasp prediction. These approaches cost very different amounts to set up, but their reported accuracies are hard to compare when the input, output, or scoring rule changes with the method. I treat the comparison as a data analysis problem. The data, the output format, and the scorer stay fixed, and the one system that gives different answers on the same input is sampled several times. The statistics are computed over objects, since objects are the independent unit here.

I built one system of each kind and kept those parts fixed. Each system saw the same 123 held-out RGB images, returned the same geometric representation, and was scored against the same Cornell labels by one implementation of the standard rectangle test [1], [2]. I also reconstructed object identities so that no physical object crossed between training and test. The statistical analysis resamples whole objects because several test images are rotated views of the same item.

I queried the VLM 615 times, five separate calls per image, so its variation on identical input could be measured. I first ask which system predicts the most valid rectangles, then look at where each system goes wrong, and finally test the most natural explanation of the VLM’s failure with a fourth run that removes its coordinate output. The three systems missed different parts of the metric. I report those failures with a common taxonomy and a separate check of whether each predicted center lands near any labeled grasp region.

Current dense grasp detectors already report higher Cornell scores [4]. I am not trying to beat them. I want to compare three practical starting points on one test that does not change with the method.

2. Related Work

Jiang *et al.* introduced the Cornell dataset and its rectangle representation of a grasp [1]. The rectangle stores the center, orientation, opening, and jaw depth. Under the

TABLE 1. OBJECT-WISE CORNELL SPLIT.

Partition	Objects	Images
Train	164	620
Validation	35	140
Test	35	123
Total	234	883

usual Cornell test, a prediction passes when it is within 30° of a labeled grasp and its intersection over union (IoU) is greater than 25% [2]. Redmon and Angelova used a neural network to regress one rectangle and reported 84.9% object-wise accuracy with RGB-D input [3]. GR-ConvNet and other pixel-wise models report higher scores, although they make a dense prediction instead of regressing one rectangle for the whole image [4].

VLM-based grasping systems usually do not ask the language model for the final pixel coordinates. Set-of-Mark prompting has the model choose among regions already marked on an image [5]. VLAD-Grasp has a VLM draw the grasp, then recovers a pose with depth prediction, segmentation, and point-cloud alignment [6]. System C tests a smaller pipeline: show GPT-4o the image and ask it to return the coordinates itself. So its score says something about that way of asking. It does not set a limit on what VLM-based grasping can do.

3. Experimental Design

3.1. Dataset, Split, and Metric

Cornell contains 885 RGB-D images of household objects, with several positive grasp rectangles on most images. I used only RGB. Two images were dropped because their object boundaries could not be resolved, leaving 883.

Cornell does not include object IDs, so I reconstructed them from the frame sequence. The procedure segmented the object from the photography platform and compared consecutive masks with color and area descriptors that tolerate rotation. Ambiguous boundaries were reviewed manually. The two possible mistakes are not equally bad, so I set the thresholds unevenly. If two objects are merged, the combined group still stays in one split. If one object is split in two, its images can leak into both training and test.

An automated first pass assigned a decision and confidence to each ambiguous boundary. In a blind 30-boundary audit, its calls agreed with my labels on 86.7% of confident cases and 40.0% of uncertain ones. I then decided every uncertain call and every high-risk “different object” call. Table 1 shows the final seed-42 split. A second check confirmed that no object appears in more than one partition.

I scored each prediction against every positive rectangle on its image. A match to any one of them counts as correct. Before using the evaluator at scale, I checked hand-computed cases, rectangle round trips, and an independent rasterized IoU calculation. All predictions were converted

back to the native 640×480 coordinate frame for the final comparison.

3.2. System A: Detector-Assisted Heuristic

System A starts with a COCO-pretrained Faster R-CNN using a ResNet50-FPN-v2 backbone [7]. A fixed lookup table maps the detected category to a center, upper, or end grasp. I calibrated the global opening and jaw-size constants on training labels only. The COCO vocabulary misses most Cornell objects, so a platform-segmentation fallback supplies a box when the detector cannot. The rule closes the gripper across the box’s narrow dimension. It can output only 0° or 90° .

The first test run scored 40.7%. Its rendered predictions exposed a defect: some accepted detections covered background clutter instead of the object on the platform. I added a guard requiring the detection to contain the segmented object’s centroid and recalibrated the global constants on training data. The score rose to 57.7%. Test output prompted that change, so 57.7% is post-hoc and 40.7% remains the clean pre-correction result.

3.3. System B: Learned Predictor

System B replaces ResNet18’s ImageNet classifier [8] with separate heads for center, orientation, and size. The orientation head predicts $(\cos 2\theta, \sin 2\theta)$. Doubling the angle handles the fact that a parallel-jaw grasp is unchanged by a 180° rotation. For images with several labels, I compute the loss against every rectangle and backpropagate only the best match. Averaging all valid labels could put the target between two good grasps, where neither one actually exists.

I trained a custom CNN, ResNet18, and ResNet34 on the 620 training images. Affine augmentation moved the rectangle corners with the pixels. Architecture, checkpoint, size-loss weight, and the custom CNN’s learning rate were chosen on validation data. ResNet18 had the best validation accuracy at 83.6%, so I froze that model before opening the test results.

That first round trained each architecture once, at one seed, with a constant learning rate, and kept the best of roughly a hundred noisy validation epochs. A second round kept the model, loss, data, and split and changed only the recipe: cosine learning-rate decay over a fixed 100 epochs, an exponential moving average of the weights, three seeds per architecture, and no checkpoint selection, with the final averaged weights reported. Test-time augmentation over eight dihedral views, combined by taking the medoid rectangle, was adopted only because it raised mean validation accuracy. Four selection rules were written down before the test split was opened a second time, and every configuration that touched test is reported.

3.4. System C: Zero-Shot VLM

System C sends the native RGB image to GPT-4o through OpenRouter [9]. It uses no grasp-specific training.

The prompt asks for two fingertip coordinates and a jaw width in JSON, which together define the rectangle. I developed the prompt on 30 training images and froze it before test.

Each test image was sent in five independent, single-turn calls at the provider’s default temperature. I retried transport failures, but not replies that arrived and failed to parse. Retrying those would quietly turn a single-call baseline into a best-of-several system. Of 615 calls, 614 parsed. The primary score pools the five exchangeable runs to estimate one-call accuracy. Best-of-five is only a ceiling: it requires five calls and an oracle that knows when one of them is right.

3.5. System D: Marked-Candidate Selection

If GPT-4o can see where to grasp but cannot express it as coordinates, removing the coordinate output should recover its choice. System D shows the same model the same test images with candidate grasps drawn on a zoomed crop and asks for the number of the one it would use. Candidates come only from the platform-segmentation mask: principal-component analysis gives the centroid, long axis, and extent, and candidates are positions along the long axis crossed with four jaw-travel directions 45° apart, so every grasp angle is within the metric’s tolerance of some candidate. Opening and jaw width use System A’s train-calibrated constants. Numbering is re-shuffled on every call. Because the menu is fixed before the model sees it, each run carries its own floor (expected accuracy of a uniformly random choice) and ceiling (whether any candidate passes).

Three menus were run, each on the full test split with five repeats. The first had twelve marks at three positions. The second had four marks at the centroid only, drawn at real size. The third had the same four marks drawn at one uniform length, and its prompt said that length is not to scale. The last two exist because on thin objects the correct marks are the smallest on the image, so a preference for large marks would mimic a preference for the long-axis grasp. The prompt was developed on 30 training images for parse rate and glyph reading only, and each menu’s test run was called once under its own sentinel.

3.6. Statistical and Failure Analysis

The 123 test images are repeated views of 35 unseen objects, not 123 independent samples. I formed the 95% confidence intervals from 20,000 object-clustered bootstrap resamples. Every view of an object, along with all five VLM calls, moves together. Paired differences use the same resamples and 20,000 object-level label-swap permutations.

The failure analysis gives each attempt one label: correct, angle only, overlap only, both criteria, or no prediction. I also checked whether the predicted center falls inside the convex hull of all labeled grasp corners on the image. This is a loose spatial check, not a claim that the grasp would work on hardware. A center outside that hull falls beyond the area spanned by the annotated grasp corners.

TABLE 2. ACCURACY WITH OBJECT-CLUSTERED 95% CONFIDENCE INTERVALS.

System	Accuracy	95% CI
A: heuristic, corrected	57.7%	[46.2, 68.2]
B: ResNet18, round 1	79.7%	[70.8, 87.1]
B: ResNet34, round 2, 3 seeds	84.0%	[73.9, 92.1]
C: GPT-4o, one call	12.4%	[8.0, 17.4]
C: best of five	35.0%	[24.6, 45.5]

4. Results

4.1. Overall Accuracy

ResNet18 passed 98 of 123 images (79.7%). The corrected heuristic passed 71 (57.7%), and GPT-4o passed 76 of 615 calls (12.4%). Its five runs ranged from 9.8% to 14.6%. At least one GPT-4o call passed on 43 of the 123 images, which gives a best-of-five ceiling of 35.0%. Table 2 includes the object-clustered intervals.

ResNet18 led the corrected heuristic by 22.0 points. The bootstrap interval for that difference was [9.7, 35.1], with permutation $p = 0.0031$. The corrected heuristic led GPT-4o’s best-of-five ceiling by 22.8 points [9.2, 35.6], $p = 0.0060$, and its expected one-call score by 45.4 points [33.8, 55.6], $p < 0.0001$. Giving each object equal weight produced 51.3%, 75.4%, and 10.5% for A, B, and one-call C. That calculation changed the size of the gaps but not their order.

The second training round chose ResNet34 on mean final-weight validation accuracy (81.4% against 81.2% for ResNet18) and used augmentation. Its three seeds scored 84.6, 83.7, and 83.7% on test, a mean of 84.0% with interval [73.9, 92.1]. ResNet18 under the same recipe averaged 82.6%. Re-running the first-round recipe at two new seeds gave 82.9% and 80.5%, so 79.7% was a low-side single draw, and the first round’s nine-point gap between ResNet18 and ResNet34 was run-to-run noise. Round 2 leads A by 26.3 points [14.8, 38.3], $p = 0.0002$. Its 4.3 point lead over round 1 does not exclude zero.

4.2. Failure Patterns

The error counts in Table 3 tell the systems apart better than the accuracy numbers do. Thirteen of A’s 52 failures had enough overlap and missed only the angle test, which fits a system limited to two orientations. B had four angle-only failures. Fifteen of its 25 misses failed only on overlap, so placement was the larger remaining problem. C failed both tests on 365 of 615 calls (59.3%).

The center-hull check shows where C’s failures begin. A landed outside the labeled grasp hull on 8 of 115 predictions (7.0%), and B did so on 8 of 123 (6.5%). C did so on 328 of 614 parsed calls (53.4%). Figure 1 shows the shared coordinate frame on one image. This is not an even contest by itself: A starts from a detected or segmented object, and B was trained on Cornell. C can place its coordinates anywhere.

TABLE 3. SHARED FAILURE TAXONOMY. PERCENTAGES USE ALL ATTEMPTS.

Outcome	A (n = 123)	B (n = 123)	C (n = 615)
Correct	71 (57.7)	98 (79.7)	76 (12.4)
Angle only	13 (10.6)	4 (3.3)	85 (13.8)
Overlap only	18 (14.6)	15 (12.2)	88 (14.3)
Both	13 (10.6)	6 (4.9)	365 (59.3)
No prediction	8 (6.5)	0 (0.0)	1 (0.2)




Figure 1. One test image under the shared coordinate system. Labeled grasps are green, System A red, System B blue, and System C’s five repeats orange.

I use this check only to see where C’s failures start. The three rates should not be read as measuring the same skill.

The five GPT-4o calls also disagreed with one another. Mean pairwise rectangle agreement was 22.0%, and mean pairwise IoU was 0.13. All five calls failed on 80 of 123 images. All five passed on one image. The remaining 42 images were unstable: identical pixels sometimes produced a passing grasp and sometimes did not.

4.3. System D

Table 4 gives the three menus. On none of them does the model’s one-call accuracy differ from a random choice among the same candidates (paired object-clustered differences -1.4 $[-6.6, +4.0]$, $+1.9$ $[-2.2, +6.3]$, and -2.3 $[-6.3, +1.7]$), while the ceilings are 79.7% and 65.9%. Its choices are not random, though. On the full menu it chose the candidate whose jaws travel along the object’s long axis near one end in 61% of calls, a grasp type that passes 17% of the time, with reasoning such as “grips the narrow top edge near the center of mass”. With four marks at the centroid drawn at one uniform size, it still chose the long-axis direction in 45% of calls against a 25% chance rate, and the short-axis candidate, which passes on 55% of images, in 17%. Stated confidence was “high” on 556 of 560 full-menu calls, 21% of them correct. Every reply parsed. The candidate at the centroid closing across the short axis, with no model at all, scored 55.3% $[41.8, 68.0]$, level with System A.

TABLE 4. SYSTEM D ON THREE CANDIDATE MENUS, SAME MODEL AND TEST IMAGES, FIVE REPEATS EACH. “LONG AXIS” IS THE SHARE OF CALLS CHOOSING THE ORIENTATION WHOSE JAWS TRAVEL ALONG THE OBJECT’S LONG AXIS; CHANCE IS 25% ON THE FOUR-MARK MENUS.

Menu	Marks	Acc.	Floor	Ceil.	Long axis
Full, 3 positions	12	18.7	20.1	79.7	70%
Centroid, real	4	26.5	24.6	65.9	41%
Centroid, uniform	4	22.3	24.6	65.9	45%

5. Discussion and Limitations

On this setup, asking GPT-4o for raw coordinates worked poorly. I am not claiming that about VLM grasping methods in general. Set-of-Mark and VLAD-Grasp use a VLM without making it emit the final coordinates, and both add machinery that System C does not have [5], [6].

On a sample of C’s failures, I read rationales that named a plausible part of the object while the coordinates landed elsewhere, which together with the 53.4% center-hull rate suggested a coordinate-binding hypothesis: the model recognizes a useful region in words and loses it while mapping the answer into pixels. System D tested that hypothesis and it failed. Taking the coordinates away and offering a menu with a 79.7% ceiling did not recover the choice, and the preference for closing along the long axis survived menus with no end-straddling candidates and with every mark the same size. From a top-down photograph, two fingertip pads straddling the end of an elongated object do look like they pinch a narrow part. In three dimensions the second finger lands on top of the object. The failure is better described as grasp reasoning from a single 2-D view than as text failing to bind to coordinates.

The experiment has several limits. It uses one object-wise split of a small benchmark, RGB without depth, and no robot trials. Cornell’s annotations are not exhaustive, so a rectangle that fails the metric might still work on a real object. System B also received more development effort: three architectures, validation-based tuning, and a second training round, compared with one frozen prompt for C and D. The test split was opened twice for System B, once per round, under rules fixed beforehand, but two openings are still two. System D used one glyph style and one candidate generator. An early formatting check on about 550 grasp rectangles probably included a few images that later entered test, although no parameter or rule came from that check. Test inspection also prompted System A’s correction. Its 40.7% score is the clean pre-correction result, and comparisons using 57.7% are post-hoc. To confirm these results I would need a fresh split or a second dataset.

6. Conclusion

ResNet18 scored highest at 79.7%, and 84.0% averaged over three seeds under a stable recipe. The corrected heuristic reached 57.7%, and GPT-4o reached 12.4% per call. Where each one goes wrong matters as much as that ranking. The heuristic often found the object but could not represent

its angle. ResNet18 usually got the angle and missed the placement. GPT-4o often failed before either distinction mattered because its rectangle was not near a labeled grasp region.

Given marked candidate grasps instead of a coordinate output, the same model chose no better than chance on three menus and kept preferring an end pinch along the object’s long axis even when every mark was drawn the same size. The problem is the grasp choice, not its expression. A label-free rule that closes across the segmented object’s short axis matches the detector-based heuristic while handling diagonal objects, and is the cheapest system here to build. What remains untested is whether depth input, or a model with a stronger three-dimensional prior, changes the choice.

Acknowledgment

OpenAI Codex helped restructure, shorten, and edit an earlier conference version of this manuscript from my existing report, code, and verified results. Anthropic Claude Code helped with code development, made the first-pass calls on ambiguous object boundaries, implemented the second training round and System D under my design, and helped edit this version. I designed the study, made the research decisions, ran and checked the experiments, verified the citations and numerical claims, and approved the final text.

GPT-4o was evaluated separately as System C in Section III-C.

References

- [1] Y. Jiang, S. Moseson, and A. Saxena, “Efficient grasping from RGBD images: Learning using a new rectangle representation,” in *Proc. IEEE Int. Conf. Robot. Autom. (ICRA)*, 2011, pp. 3304–3311.
- [2] I. Lenz, H. Lee, and A. Saxena, “Deep learning for detecting robotic grasps,” *Int. J. Robot. Res.*, vol. 34, nos. 4–5, pp. 705–724, 2015.
- [3] J. Redmon and A. Angelova, “Real-time grasp detection using convolutional neural networks,” in *Proc. IEEE Int. Conf. Robot. Autom. (ICRA)*, 2015, pp. 1316–1322.
- [4] S. Kumra, S. Joshi, and F. Sahin, “Antipodal robotic grasping using generative residual convolutional neural network,” in *Proc. IEEE/RSJ Int. Conf. Intell. Robots Syst. (IROS)*, 2020, pp. 9626–9633.
- [5] J. Yang *et al.*, “Set-of-Mark prompting unleashes extraordinary visual grounding in GPT-4V,” *arXiv:2310.11441*, 2023.
- [6] M. Kulshrestha *et al.*, “VLAD-Grasp: Zero-shot grasp detection via vision-language models,” *arXiv:2511.05791*, 2025.
- [7] S. Ren, K. He, R. Girshick, and J. Sun, “Faster R-CNN: Towards real-time object detection with region proposal networks,” in *Adv. Neural Inf. Process. Syst.*, vol. 28, 2015.
- [8] K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2016, pp. 770–778.
- [9] OpenAI, “GPT-4o system card,” *arXiv:2410.21276*, 2024.